



StephenMuuo /
Phase1-Project



<> Code

Issues

Pull requests

Actions

Projects

Wiki

Security



Phase1-Project / Jupyter Notebook.ipynb



StephenMuuo Add files via upload

66910bf · now



1256 lines (1256 loc) · 94.5 KB

Preview

Code

Blame

Raw



Project Overview: Assessing Aircraft Risk for New Aviation Division

Business problem

Your company is expanding in to new industries to diversify its portfolio. Specifically, they are interested in purchasing and operating airplanes for commercial and private enterprises, but do not know anything about the potential risks of aircraft. You are charged with determining which aircraft are the lowest risk for the company to start this new business endeavor. You must then translate your findings into actionable insights that the head of the new aviation division can use to help decide which aircraft to purchase.

The Data

The data is from the National Transportation Safety Board that includes aviation accident data from 1962 to 2023 about civil aviation accidents and selected incidents in the United States and international waters.

Goals

To identify which aircrafts are the lowest risk for the company.

Technical Presentation.

```
In [5]: #Importing Libraries.
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [6]: #Importing the dataset
df = pd.read_csv('AviationData.csv', encoding='latin-1')
```

```
C:\Users\ACDP-KENYA\anaconda3\envs\learn-env\lib\site-packages\IPython\core
\interactiveshell.py:3145: DtypeWarning: Columns (6,7,28) have mixed types.
Specify dtype option on import or set low_memory=False.
  has_raised = await self.run_ast_nodes(code_ast.body, cell_name,
```

Data Exploration.


```
In [7]: #Explore the first rows
df.head()
```

```
Out[7]:
```

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location
--	----------	--------------------	-----------------	------------	----------

0	20001218X45444	Accident	SEA87LA080	1948-10-24	MOOSE CREEK, ID
1	20001218X45447	Accident	LAX94LA336	1962-07-19	BRIDGEPORT CA
2	20061025X01555	Accident	NYC07LA005	1974-08-30	Saltville, VA
3	20001218X45448	Accident	LAX96LA321	1977-06-19	EUREKA, CA
4	20041105X01764	Accident	CHI79FA064	1979-08-02	Canton, OH

5 rows × 31 columns



In [8]: *#Checking its shape*
df.shape

Out[8]: (88889, 31)

In [9]: *#Checking the columns*
df.columns

Out[9]: Index(['Event.Id', 'Investigation.Type', 'Accident.Number', 'Event.Date', 'Location', 'Country', 'Latitude', 'Longitude', 'Airport.Code', 'Airport.Name', 'Injury.Severity', 'Aircraft.damage', 'Aircraft.Category', 'Registration.Number', 'Make', 'Model', 'Amateur.Built', 'Number.ofEngines', 'Engine.Type', 'FAR.Descripti on', 'Schedule', 'Purpose.of.flight', 'Air.carrier', 'Total.Fatal.Injuri es', 'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjure d', 'Weather.Condition', 'Broad.phase.of.flight', 'Report.Status', 'Publication.Date'], dtype='object')

In [10]: df.info

Out[10]: <bound method DataFrame.info of
accident.Number Event.Date \

0	20001218X45444	Accident	SEA87LA080	1948-10-24
1	20001218X45447	Accident	LAX94LA336	1962-07-19
2	20061025X01555	Accident	NYC07LA005	1974-08-30
3	20001218X45448	Accident	LAX96LA321	1977-06-19
4	20041105X01764	Accident	CHI79FA064	1979-08-02
...
88884	20221227106491	Accident	ERA23LA093	2022-12-26
88885	20221227106494	Accident	ERA23LA095	2022-12-26
88886	20221227106497	Accident	WPR23LA075	2022-12-26
88887	20221227106498	Accident	WPR23LA076	2022-12-26
88888	20221230106513	Accident	ERA23LA097	2022-12-29

Location Country Latitude Longitude Airport.Code \

0	MOOSE CREEK, ID	United States	NaN	NaN	NaN
1	BRIDGEPORT, CA	United States	NaN	NaN	NaN
2	Saltville, VA	United States	36.0222	81.8781	NaN

```

2         SALISVILLE, VA United States 30.7222 -01.0701 NaN
3         EUREKA, CA United States NaN NaN NaN
4         Canton, OH United States NaN NaN NaN
...
88884     Annapolis, MD United States NaN NaN NaN
88885     Hampton, NH United States NaN NaN NaN
88886     Payson, AZ United States 341525N 1112021W PAN
88887     Morgan, UT United States NaN NaN NaN
88888     Athens, GA United States NaN NaN NaN

```

```

      Airport.Name ... Purpose.of.flight Air.carrier \
0         NaN ... Personal NaN
1         NaN ... Personal NaN
2         NaN ... Personal NaN
3         NaN ... Personal NaN
4         NaN ... Personal NaN
...
88884     NaN ... Personal NaN
88885     NaN ... NaN NaN
88886     PAYSON ... Personal NaN
88887     NaN ... Personal MC CESSNA 210N LLC
88888     NaN ... Personal NaN

```

```

      Total.Fatal.Injuries Total.Serious.Injuries Total.Minor.Injuries \
0         2.0 0.0 0.0
1         4.0 0.0 0.0
2         3.0 NaN NaN
3         2.0 0.0 0.0
4         1.0 2.0 NaN
...
88884     0.0 1.0 0.0
88885     0.0 0.0 0.0
88886     0.0 0.0 0.0
88887     0.0 0.0 0.0
88888     0.0 1.0 0.0

```

```

      Total.Uninjured Weather.Condition Broad.phase.of.flight \
0         0.0 UNK Cruise
1         0.0 UNK Unknown
2         NaN IMC Cruise
3         0.0 IMC Cruise
4         0.0 VMC Approach
...
88884     0.0 NaN NaN
88885     0.0 NaN NaN
88886     1.0 VMC NaN
88887     0.0 NaN NaN
88888     1.0 NaN NaN

```

```

      Report.Status Publication.Date
0     Probable Cause NaN
1     Probable Cause 19-09-1996
2     Probable Cause 26-02-2007
3     Probable Cause 12-09-2000
4     Probable Cause 16-04-1980
...
88884     NaN 29-12-2022
88885     NaN NaN
88886     NaN 27-12-2022
88887     NaN NaN
88888     NaN 30-12-2022

```

```
[88889 rows x 31 columns]>
```

```
In [11]: df.describe
```

```
df.describe
```

```
Out[11]: <bound method NDFrame.describe of
Accident.Number  Event.Date  \
0      20001218X45444      Accident      SEA87LA080      1948-10-24
1      20001218X45447      Accident      LAX94LA336      1962-07-19
2      20061025X01555      Accident      NYC07LA005      1974-08-30
3      20001218X45448      Accident      LAX96LA321      1977-06-19
4      20041105X01764      Accident      CHI79FA064      1979-08-02
...
88884      20221227106491      Accident      ERA23LA093      2022-12-26
88885      20221227106494      Accident      ERA23LA095      2022-12-26
88886      20221227106497      Accident      WPR23LA075      2022-12-26
88887      20221227106498      Accident      WPR23LA076      2022-12-26
88888      20221230106513      Accident      ERA23LA097      2022-12-29

      Location      Country  Latitude  Longitude  Airport.Code  \
0      MOOSE CREEK, ID      United States      NaN      NaN      NaN
1      BRIDGEPORT, CA      United States      NaN      NaN      NaN
2      Saltville, VA      United States      36.9222      -81.8781      NaN
3      EUREKA, CA      United States      NaN      NaN      NaN
4      Canton, OH      United States      NaN      NaN      NaN
...
88884      Annapolis, MD      United States      NaN      NaN      NaN
88885      Hampton, NH      United States      NaN      NaN      NaN
88886      Payson, AZ      United States      341525N      1112021W      PAN
88887      Morgan, UT      United States      NaN      NaN      NaN
88888      Athens, GA      United States      NaN      NaN      NaN

      Airport.Name  ...  Purpose.of.flight      Air.carrier  \
0      NaN      ...      Personal      NaN
1      NaN      ...      Personal      NaN
2      NaN      ...      Personal      NaN
3      NaN      ...      Personal      NaN
4      NaN      ...      Personal      NaN
...
88884      NaN      ...      Personal      NaN
88885      NaN      ...      NaN      NaN
88886      PAYSON      ...      Personal      NaN
88887      NaN      ...      Personal      MC CESSNA 210N LLC
88888      NaN      ...      Personal      NaN

      Total.Fatal.Injuries  Total.Serious.Injuries  Total.Minor.Injuries  \
0      2.0      0.0      0.0
1      4.0      0.0      0.0
2      3.0      NaN      NaN
3      2.0      0.0      0.0
4      1.0      2.0      NaN
...
88884      0.0      1.0      0.0
88885      0.0      0.0      0.0
88886      0.0      0.0      0.0
88887      0.0      0.0      0.0
88888      0.0      1.0      0.0

      Total.Uninjured  Weather.Condition  Broad.phase.of.flight  \
0      0.0      UNK      Cruise
1      0.0      UNK      Unknown
2      NaN      IMC      Cruise
3      0.0      IMC      Cruise
4      0.0      VMC      Approach
...
88884      0.0      NaN      NaN
88885      0.0      NaN      NaN
88886      0.0      NaN      NaN
```

```

00000      1.0      NaN
88887      0.0      NaN
88888      1.0      NaN

```

```

      Report.Status Publication.Date
0      Probable Cause      NaN
1      Probable Cause    19-09-1996
2      Probable Cause    26-02-2007
3      Probable Cause    12-09-2000
4      Probable Cause    16-04-1980
...
88884      NaN    29-12-2022
88885      NaN      NaN
88886      NaN    27-12-2022
88887      NaN      NaN
88888      NaN    30-12-2022

```

```
[88889 rows x 31 columns]>
```

Data Cleaning

From the above findings we can see the dataset has some null values. We are going to clean the data inorder to get a better understanding of the data.

```
In [12]: #Checking for Null values
df.isnull().sum()
```

```
Out[12]: Event.Id      0
Investigation.Type    0
Accident.Number      0
Event.Date           0
Location            52
Country             226
Latitude           54507
Longitude           54516
Airport.Code        38640
Airport.Name        36099
Injury.Severity     1000
Aircraft.damage     3194
Aircraft.Category   56602
Registration.Number  1317
Make                63
Model               92
Amateur.Built       102
Number.of.Engines   6084
Engine.Type         7077
FAR.Description     56866
Schedule            76307
Purpose.of.flight   6192
Air.carrier         72241
Total.Fatal.Injuries 11401
Total.Serious.Injuries 12510
Total.Minor.Injuries 11933
Total.Uninjured     5912
Weather.Condition   4492
Broad.phase.of.flight 27165
Report.Status       6381
Publication.Date    13771
dtype: int64
```

```
In [13]: #Drop columns that have more missing values
threshold = len(df) * 0.5
```

```
df = df.dropna(thresh=thershold, axis=1)
df.shape
```

Out[13]: (88889, 25)

```
In [14]: #Finding more Null Values
df.isnull().sum()
```

Out[14]:

Event.Id	0
Investigation.Type	0
Accident.Number	0
Event.Date	0
Location	52
Country	226
Airport.Code	38640
Airport.Name	36099
Injury.Severity	1000
Aircraft.damage	3194
Registration.Number	1317
Make	63
Model	92
Amateur.Built	102
Number.of.Engines	6084
Engine.Type	7077
Purpose.of.flight	6192
Total.Fatal.Injuries	11401
Total.Serious.Injuries	12510
Total.Minor.Injuries	11933
Total.Uninjured	5912
Weather.Condition	4492
Broad.phase.of.flight	27165
Report.Status	6381
Publication.Date	13771
dtype:	int64

```
In [15]: #Checking the difference of the remaining values
d={}
for i in df.columns:
    d[i]=len(df[i].unique())
unique_values= pd.DataFrame(list(d.items()), columns=['Column', 'unique_val'])
unique_values["missing_values"]= df.isna().sum().values
unique_values
```

Out[15]:

	Column	unique_val	missing_values
--	--------	------------	----------------

0	Event.Id	87951	0
1	Investigation.Type	2	0
2	Accident.Number	88863	0
3	Event.Date	14782	0
4	Location	27759	52
5	Country	220	226
6	Airport.Code	10376	38640
7	Airport.Name	24872	36099

8	Injury.Severity	110	1000
9	Aircraft.damage	5	3194
10	Registration.Number	79106	1317
11	Make	8238	63
12	Model	12319	92
13	Amateur.Built	3	102
14	Number.of.Engines	8	6084
15	Engine.Type	14	7077
16	Purpose.of.flight	27	6192
17	Total.Fatal.Injuries	126	11401
18	Total.Serious.Injuries	51	12510
19	Total.Minor.Injuries	58	11933
20	Total.Uninjured	380	5912
21	Weather.Condition	5	4492
22	Broad.phase.of.flight	13	27165
23	Report.Status	17076	6381
24	Publication.Date	2925	13771

In [16]:

```
#Dropping columns with more missing values
subset_col=list(unique_values[unique_values['missing_values']>1000]['Column'])
df.dropna(subset=subset_col, inplace=True)
df.shape
```

<ipython-input-16-518fa402ec84>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df.dropna(subset=subset_col, inplace=True)

Out[16]: (18700, 25)

In [17]:

```
df.isnull().sum()
```

```
Out[17]: Event.Id          0
Investigation.Type       0
Accident.Number         0
Event.Date              0
Location                0
Country                 47
Airport.Code            0
Airport.Name            0
Injury.Severity         0
Aircraft.damage         0
Registration.Number     0
Make                    3
.. ..
```



```

model          12
Amateur.Built   0
Number.of.Engines  0
Engine.Type     0
Purpose.of.flight  0
Total.Fatal.Injuries  0
Total.Serious.Injuries  0
Total.Minor.Injuries  0
Total.Uninjured  0
Weather.Condition  0
Broad.phase.of.flight  0
Report.Status    0
Publication.Date  0
dtype: int64

```

In [18]:

```

#Filling in Null values
df.fillna(method='ffill', inplace=True)
df.isnull().sum()

```

C:\Users\ACDP-KENYA\anaconda3\envs\learn-env\lib\site-packages\pandas\core\frame.py:4317: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
 return super().fillna()

Out[18]:

```

Event.Id          0
Investigation.Type  0
Accident.Number   0
Event.Date        0
Location          0
Country           0
Airport.Code      0
Airport.Name      0
Injury.Severity   0
Aircraft.damage   0
Registration.Number  0
Make             0
Model            0
Amateur.Built     0
Number.of.Engines  0
Engine.Type       0
Purpose.of.flight  0
Total.Fatal.Injuries  0
Total.Serious.Injuries  0
Total.Minor.Injuries  0
Total.Uninjured   0
Weather.Condition  0
Broad.phase.of.flight  0
Report.Status     0
Publication.Date   0
dtype: int64

```

In [19]:

```

#Checking for Duplicates
df.duplicated().sum()

```

Out[19]: 0

I have cleaned the data set, there may be some errors but I believe they will not cause an impact on the final findings of the investigation.

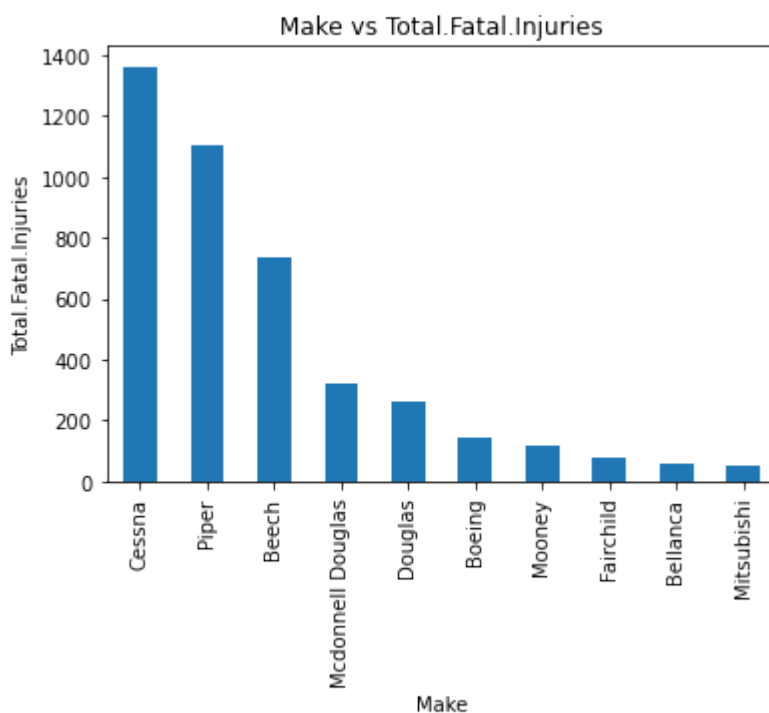
Data Visualisation and Exploratory Data Analysis.

From the above cleaning and exploration we can visualise some findings.

Identifying the riskiest Make against Fatal Injuries.

In [20]:

```
#Visualising Make vs Total.Fatal.Injuries
df.groupby('Make')['Total.Fatal.Injuries'].sum().sort_values(ascending=False)
plt.title('Make vs Total.Fatal.Injuries')
plt.xlabel('Make')
plt.ylabel('Total.Fatal.Injuries')
plt.show()
```




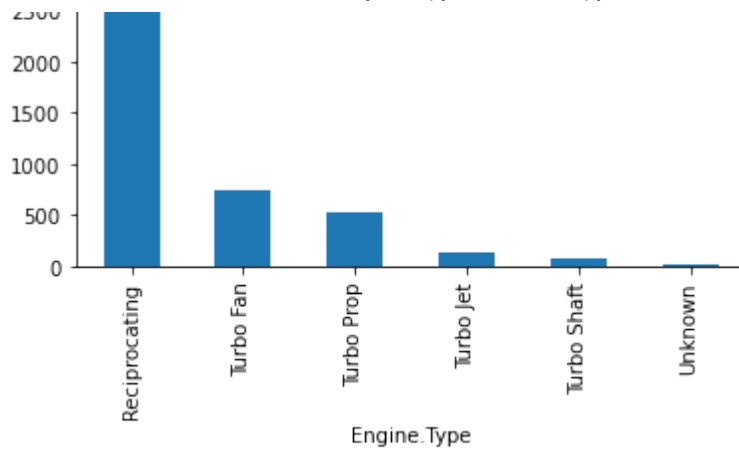
Identifying the riskiest Engine type against total fatalities.

In [21]:

```
#Visualisation of Engine.Type vs Total.Fatal.Injuries.
df.groupby('Engine.Type')['Total.Fatal.Injuries'].sum().sort_values(ascending=False)
plt.title('Engine.Type vs Total.Fatal.Injuries')
plt.xlabel('Engine.Type')
plt.show()
```

Out[21]: <module 'matplotlib.pyplot' from 'C:\\Users\\ACDP-KENYA\\anaconda3\\envs\\learn-env\\lib\\site-packages\\matplotlib\\pyplot.py'>






Identify which purpose of flight had the highest risk.

In [22]:

```
#Visualisation of Purpose.of.flight vs Total.Fatal.Injuries.  
df.groupby('Purpose.of.flight')['Total.Fatal.Injuries'].sum().sort_values()  
plt.title('Purpose.of.flight vs Total.Fatal.Injuries')  
plt.xlabel('Purpose.of.flight')  
plt.ylabel('Total.Fatal.Injuries')  
plt.show()
```



Recommendations.

Aircraft Make

Cesena and Piper are the highest risk aircrafts and should be avoided. Mitsubishi had the lowest fatalities.

Aircraft Engine

Reciprocating Engine had the highest fatalities and should be avoided. Turbo shaft

engines had very low fatalities and should be considered

Purpose of Flight

Personal perpose of flights had the highest fatalities and carry the highest risk of reason for flight.

Next Steps.